

NoSQL concepts



Data Analytics Tech Lead & Product Manager Worldline

2013 - 2016

Big Data
Engineer

2016 - 2017

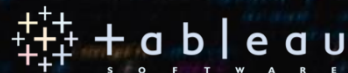
Data Analytics
Evangelist

2017-2019

Data Science
Tech Lead

2019-current

Data Analytics
Product Manager

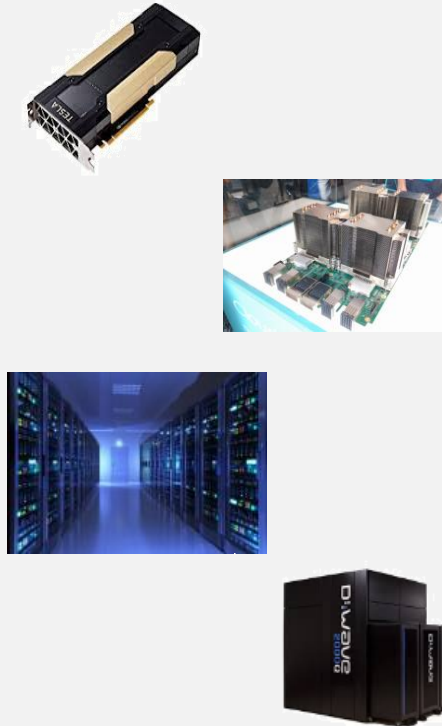


Data is at the center of all IT activities

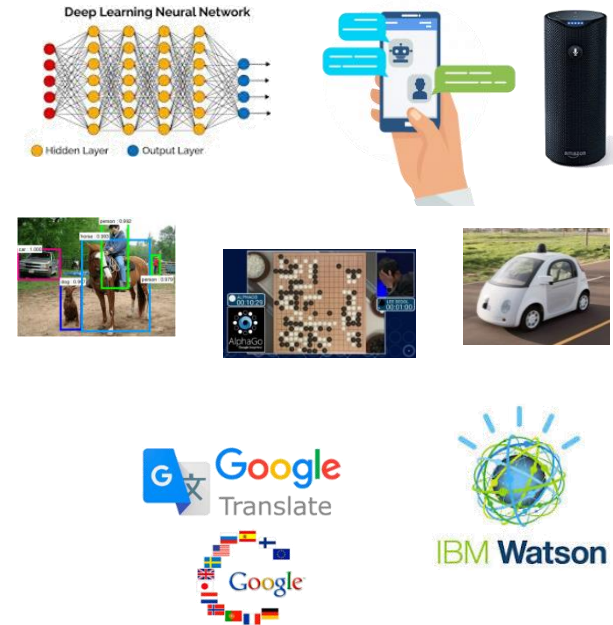
Data



Hardware



Innovations



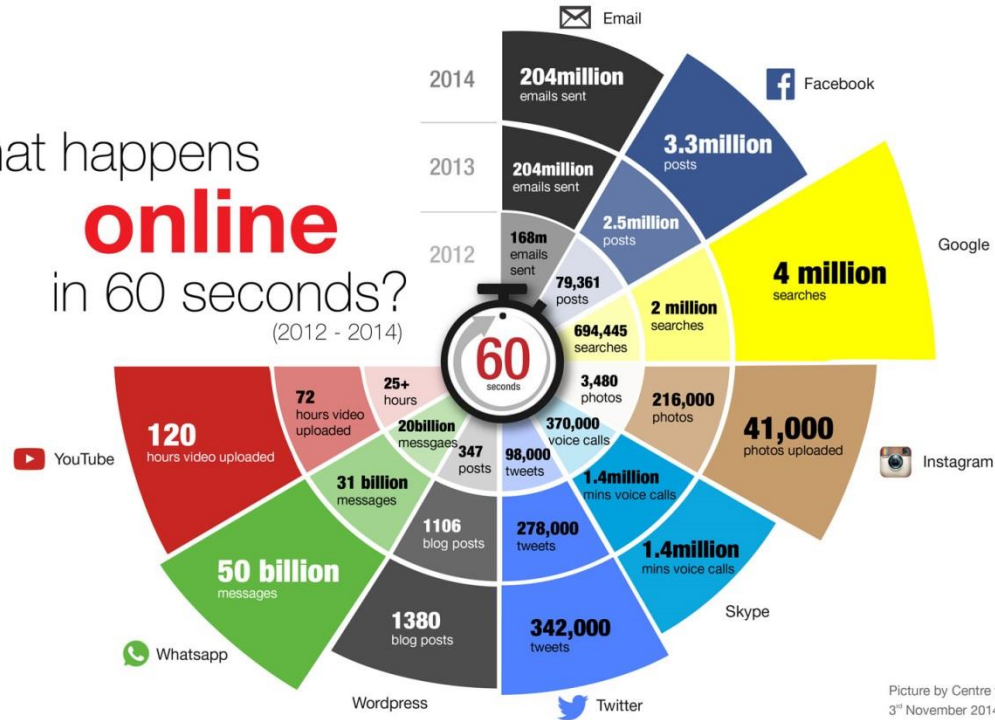


Data Scientist:

The Sexiest Job of the 21st Century

Problem

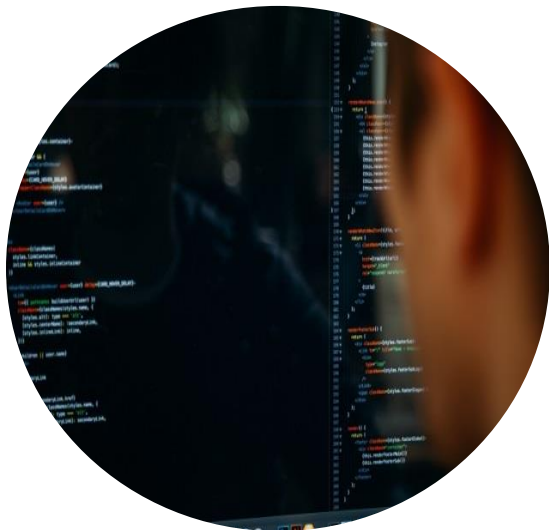
What happens
online
in 60 seconds?
(2012 - 2014)



Picture by Centre for Learning and Teaching
3rd November 2014

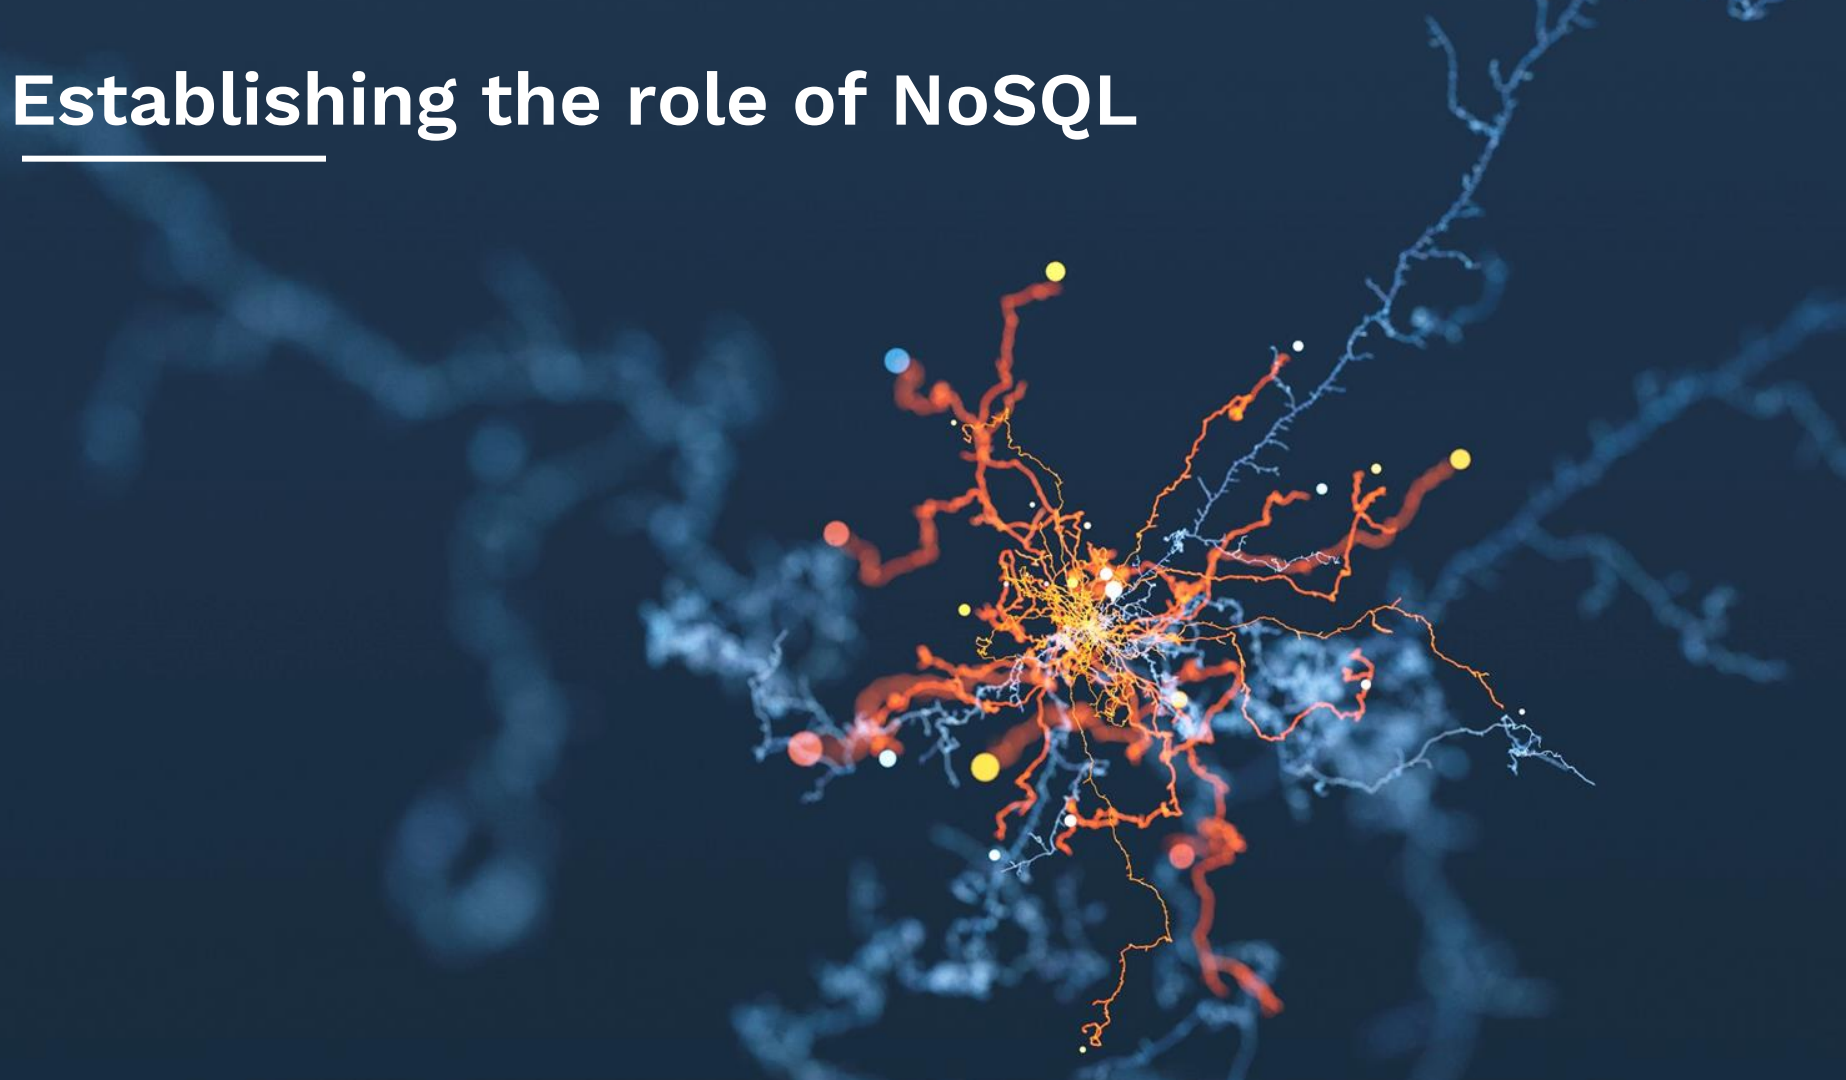
How should we store and query such data ?

Agenda



- ❑ Establishing the role of NoSQL
- ❑ NoSQL data stores
- ❑ CAP theorem
- ❑ Denormalization
- ❑ Hands-on

Establishing the role of NoSQL

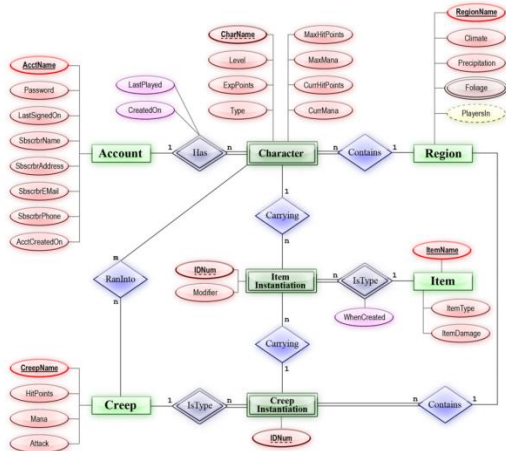




Monopoly of RDBMS since 1970

- 1970: Edgar F. Codd theorizes the relational model
- 1974: Creation of the SQL language
- 1978: Multics, 1st commercial solution
- 1979: 1st Oracle solution
- 1986: SQL is standardized
- 2000: MySQL is open-sourced

RDBMS



Relational paradigm

- Database = set of tables and relations
- Table = set of tuples (rows) structured under columns
- Column = schema constraint on the data

Perfect for business informatics

The emergence of Not Only SQL

Big Data and the Web

- ▶ Velocity, Variety, Volume
- ▶ Simple storage problems



... limited by their tools

- ▶ Vertical scalability
- ▶ Relational paradigm is rigid
- ▶ Expensive !



Best quality of service

- ▶ Aim for no service interruption
- ▶ Aim for best user experience
- ▶ Service/data in constant evolution



The emergence of Not Only SQL

Summer 2009: NoSQL Meetup in California

- ▶ Talk about distributed storage systems
- ▶ Reuse the NoSQL term from Carlos Strozzi in 1998
- ▶ Voldemort (LinkedIn), Cassandra (Facebook), Dynomite, Hbase, Hypertable, CouchDB

A definition



**Non-relational, distributed, horizontally
scalable and open-source database**



<http://nosql-database.org>

Beyond the relational model

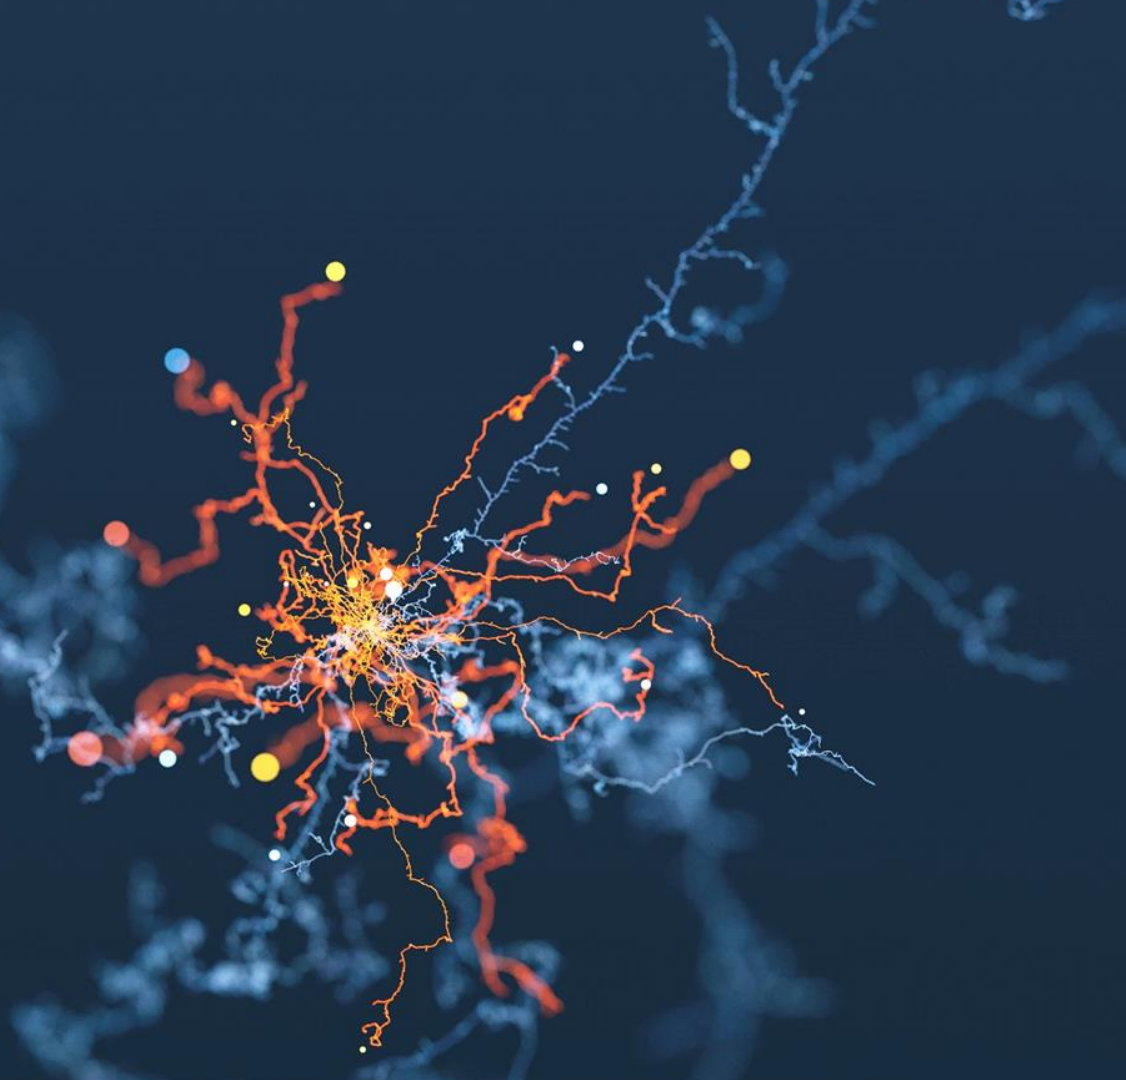


- Consistency makes it hard to handle scalability and millions of users connecting
- Relational model rigidity makes it hard to modify a model on the fly without slowing the service

Complements (and doesn't replace) SQL

- ▶ Responds to new requirements : scalability, specialization in data storage
- ▶ Different constraints & paradigms

NoSQL Data stores



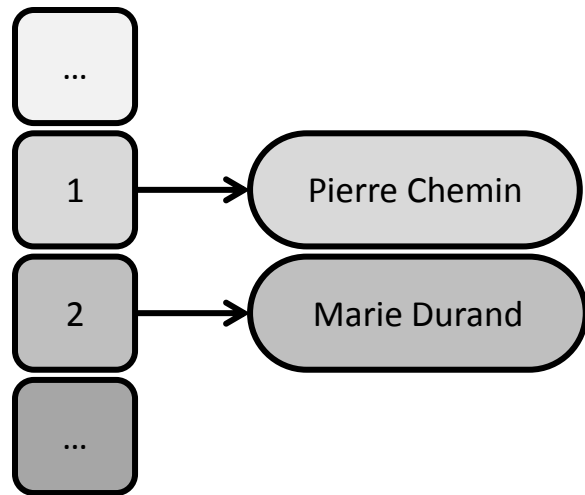
Data structures



So what data structures do you know in
Computer Science ?

Key value stores

- ▶ N keys of type String
 - ▶ For each key, an opaque value
 - ▶ CRUD operations
-
- ⊕ Simplicity, performance
 - ⊖ Cannot perform query on values
 - ⊖ Not suited for complex modelisations



Ex : Redis, Riak, Memcached, Voldemort



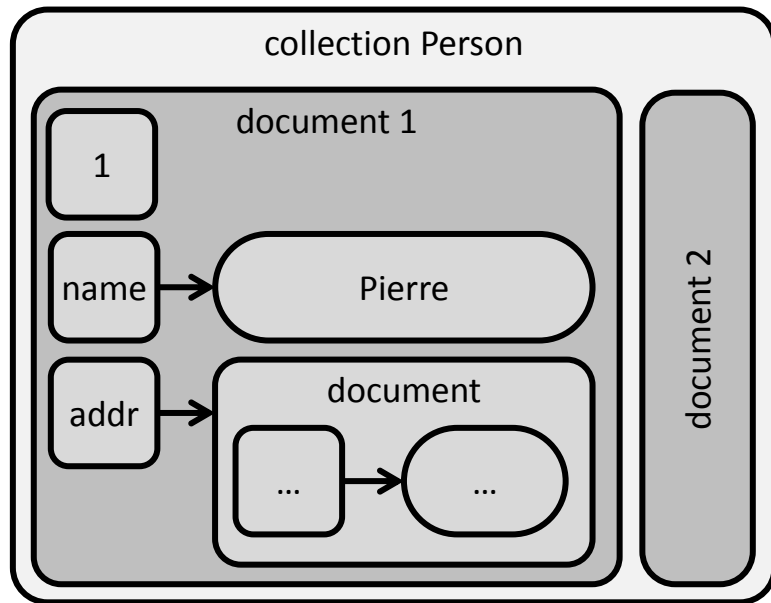
Document

- ▶ N collections of documents
 - ▶ A document = id + k fields
 - ▶ A field = name + value
 - ▶ CRUD + query language
-
- ⊕ Native integration with object language
 - ⊕ Querying on any field
 - ⊖ Need to denormalize

Ex : MongoDB, CouchDB, RethinkDB



RethinkDB

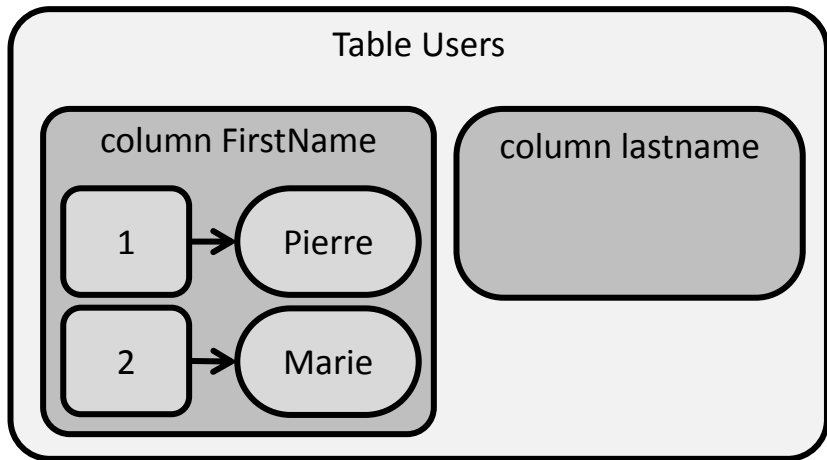


Column-oriented

- ▶ A table = N columns
- ▶ A tuple = set of columns associated to the same key
- ▶ CRUD + query per column

- ⊕ Tuple is a well known entity
- ⊕ Easy to partition and rebalance
- ⊖ Need to denormalize

Ex: Hbase, Cassandra, HyperTable

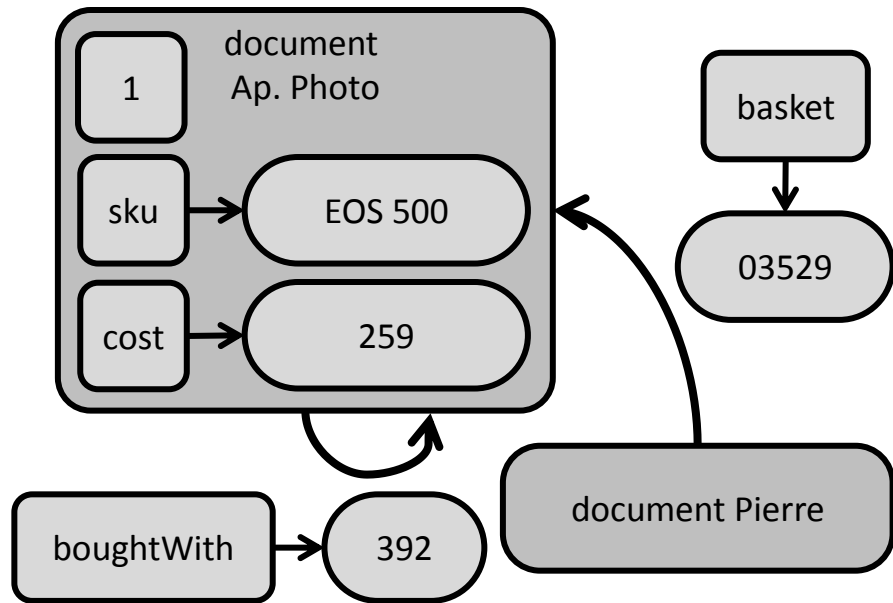


APACHE
HBASE



Graph

- ▶ Vertices linked by edges
 - ▶ Vertex = a document
 - ▶ Edge = a document with origin and destination
 - ▶ CRUD + query language
- ⊕ Modelisation close to object oriented
- ⊖ More complex usage
- ⊖ Not as scalable as other



Ex: Neo4j, OrientDB, Titan (Cassandra / Hbase)



Practice ?



- ☐ File caching (Web static resources)
- ☐ Volatile data (counters, chats, basket...)

- ☐ Analysis of transactions (eCommerce, finance...)
- ☐ Content storage (mail, messages...)

- ☐ Content Managing System (Wiki, Blog, Forum...)
- ☐ Internet of Things (Machine2Machine, sensors)

- ☐ Recommendation/suggestion (retail, social...)
- ☐ Social networks

Fanilo → 36

Adam → 54

Weather → [cold, hot]

Last_date → 09/08/2016

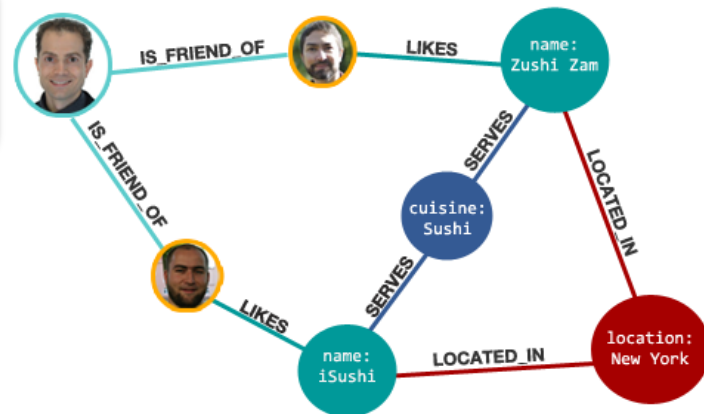
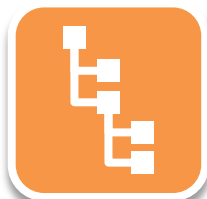


APACHE
HBASE



0	pomme		2,99	
7	poire		5,99	
3		red	12,00	music
		blue	89	

```
{ "menu":  
  { "id": "file",  
    "value": "File",  
    "popup": {  
      "menuitem": [  
        { « v": "New", "onclick": "new()" },  
        { « v": "Open", "onclick": "open()" },  
        { « v": "Close", "onclick": "close()" }  
      ]  
    }  
  }  
}
```





RELAX

CAP theorem

BRACEYOURSELF



Physical constraints on distributed systems

Distributed systems has two natural constraints
so they don't behave like a single system !



Number of nodes



Probability of the failure of the system
Need of communication between nodes



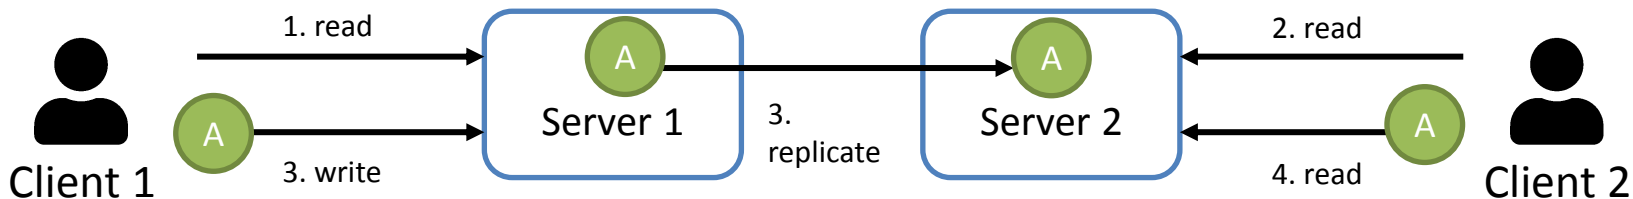
Distance between nodes



Minimum latency for communication

Strong consistency

All nodes see the same data at the same time

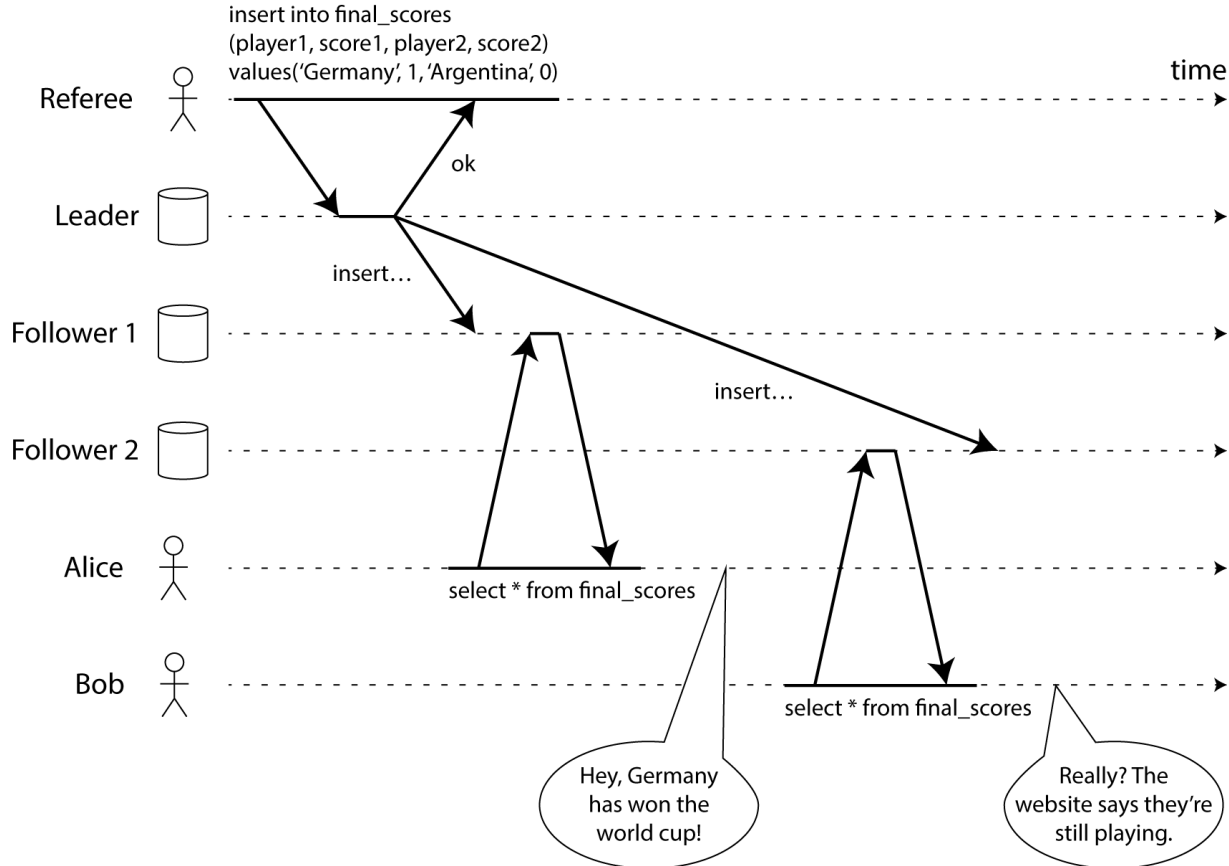


- Linearizable consistency: all operations appear to have executed atomically in an order that is consistent with the global real-time ordering of operations
- Sequential consistency: all operations appear to have executed atomically in some order that is consistent with the order seen at individual nodes and that is equal at all nodes



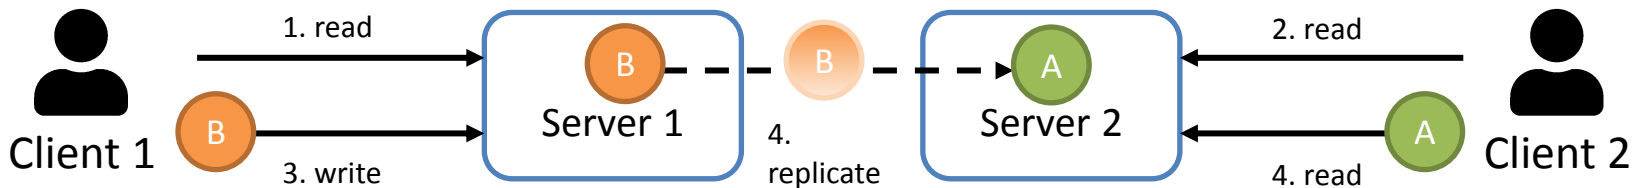
- ❑ Needs to block writing during the update
- ❑ Blocking time is unpredictable given network issues

Non-linearizable system



Weak consistency

If you stop changing values, then after some undefined amount of time all replicas will agree on the same value



Saying something is eventually consistent is like saying “people are eventually dead” :

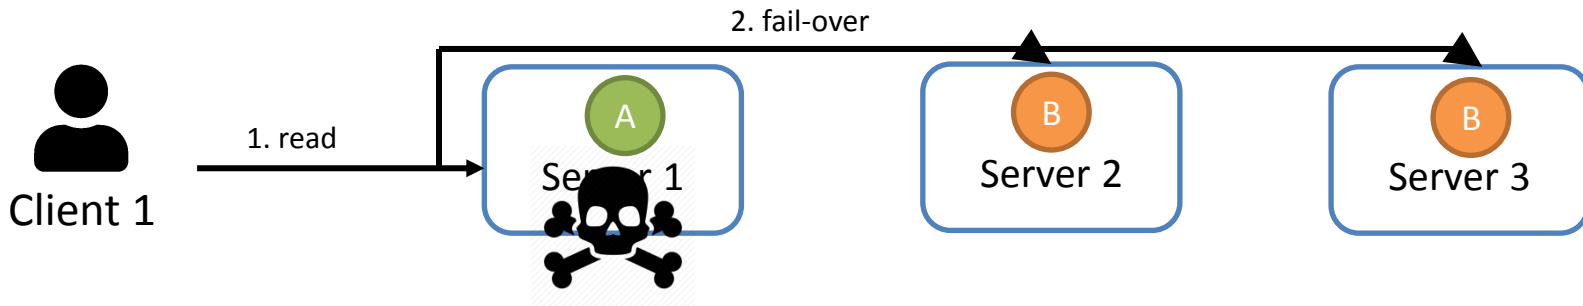
- how long is « eventually » ? (ex: Get a strict lower bound)
- how do the replicas agree on a value ? (ex: value with the largest timestamp always win)



Non-blocking writing : better performance,
at the risk of reading stale values

Availability

Node failures do not prevent survivors from continuing to operate



Use replicas:

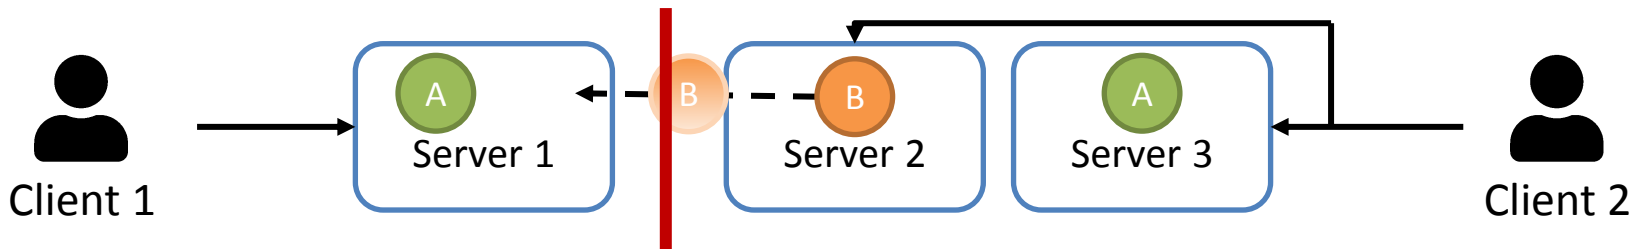
- For reading for better performance (nearest copy)
- For writing, if machine where we should write is dead (conflict resolution)

Failure tolerance

- Heartbeating / Gossip protocols to identify dead nodes
- Fail-over : when an original is dead, a replica is promoted to original, update all machines/clients to point to new original

Partition tolerance

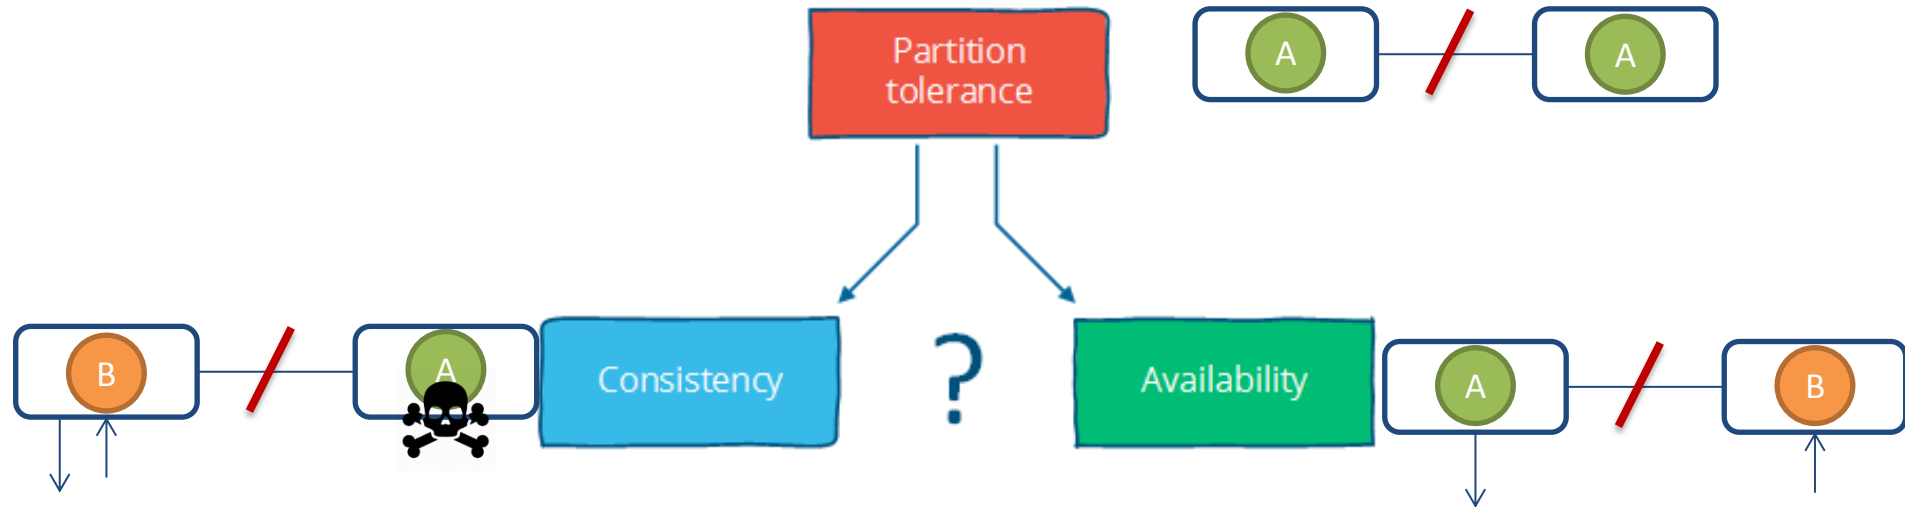
The system continues to operate despite message loss due to network and/or node failure



Divided machines must work independently

- Can still read and write data
 - When network is back, should resolve data conflicts between different machines
- Many system designs used in early distributed relational database systems do not take into account partition tolerance

CAP theorem

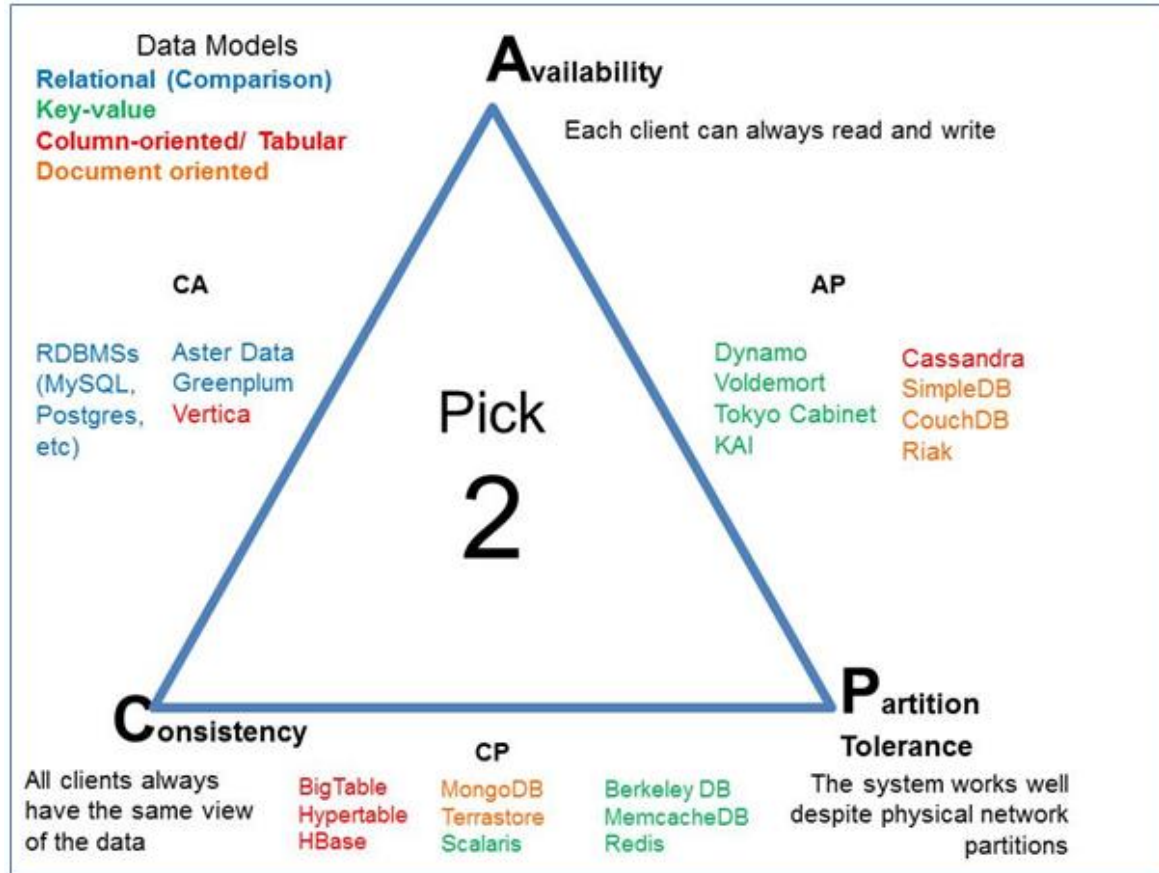


On a network failure

- Choose strong consistency → Refuse the write, knowing that the client might not be able to contact A or B until the partition heals.
- Choose availability → Accept the write, knowing that neither A nor B will know about this new data until the partition heals.

The CAP theorem is an example of distributed system \neq single server

CAP theorem





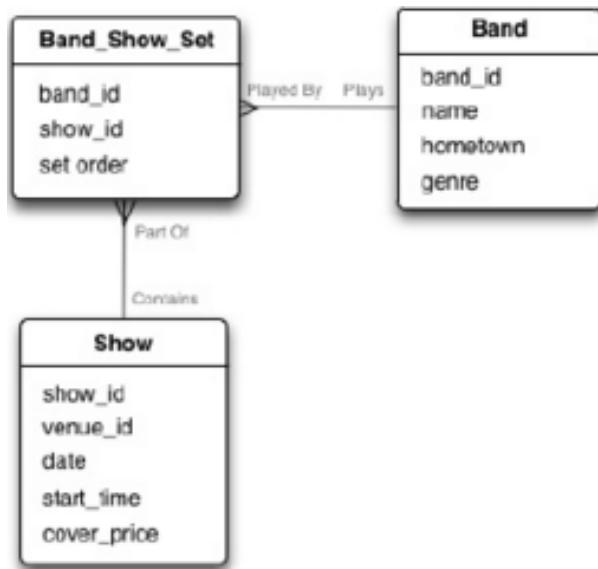
RELAX

Denormalization



Denormalization

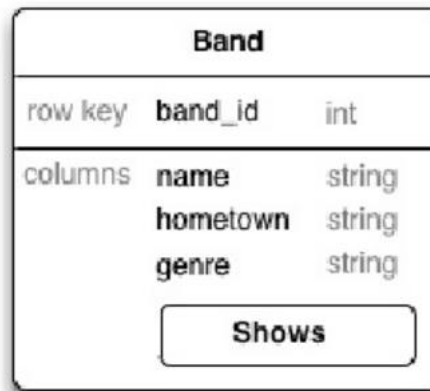
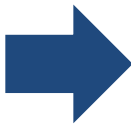
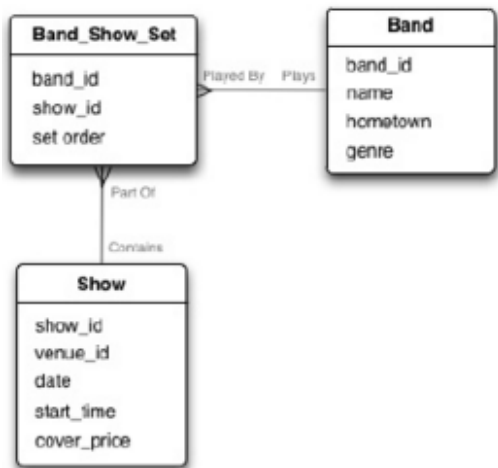
- We are used to representing relationships between entities in the relational model in a normalized way
 - In the DB we then perform joins on tables as optimally as we can



- But...the relational model doesn't acknowledge scale and we need to cheat
 - Add indexes, add hints, write complex SQL, trickle with the optimizer, cache...

Denormalization

- NoSQL is a different way of physically representing your logical data model
 - + Scalable, efficient
 - “No” foreign keys, “nor” joins, “nor” cross-table transactions
- Denormalization : *logical* entities share one *physical* representation
 - You can have one of your fields representing another entity → nested entity !



Denormalization

Band		
row key	band_id	int
columns	name	string
	hometown	string
	genre	string
Shows		
Albums		
Members		



What do you think about querying this structure ?

Denormalization

Band		
row key	band_id	int
columns	name	string
	hometown	string
	genre	string
Shows		
Albums		
Members		



- You then get all shows per band for free and fast
- You can put more nested entities
- You can nest more entities into the nested entity
- In the end, you create one denormalized table per business query you need



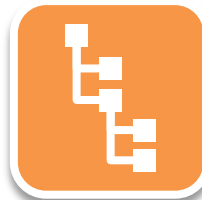
-
- Nested entities aren't independent anymore
 - "SELECT avg(price) FROM albums WHERE name = "bla"
 - Decomposition is unilateral
 - Makes you think more about physical access than relational model



CONCLUSION

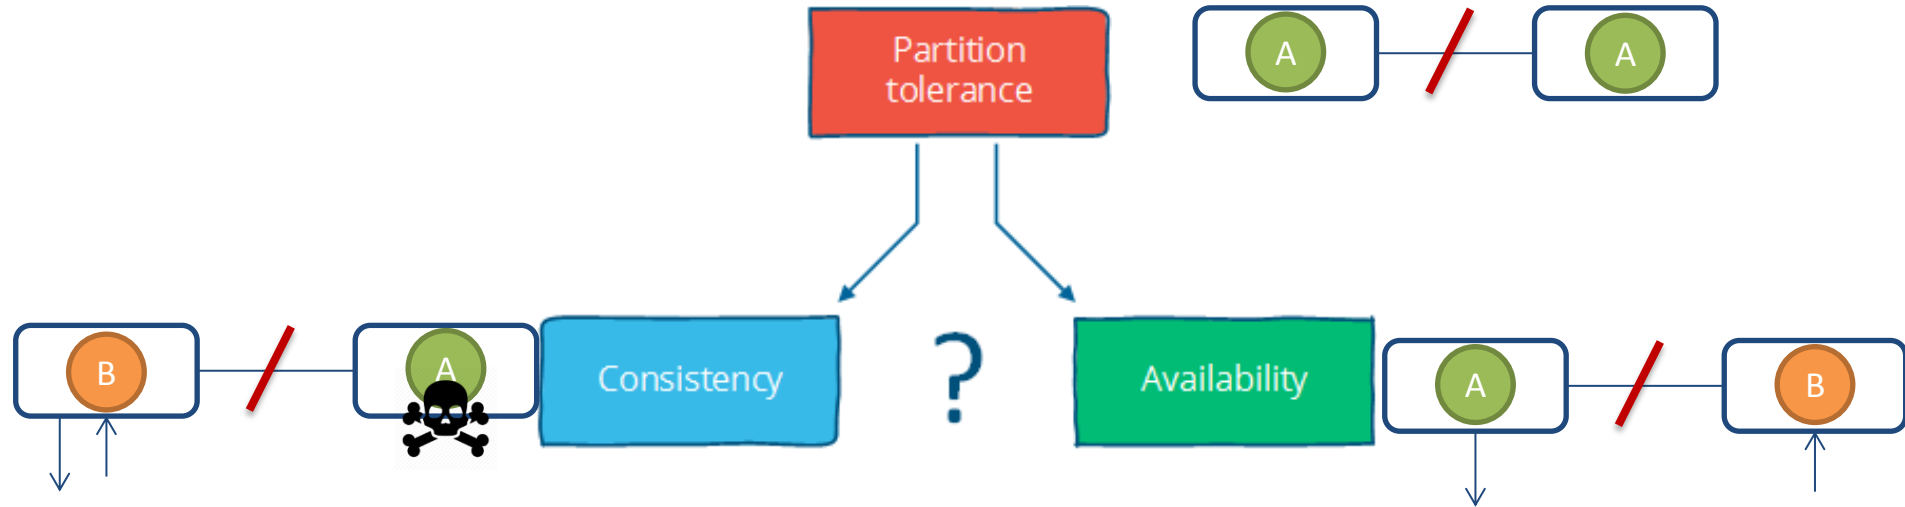
Conclusion

- A **NoSQL** (originally referring to "non SQL" or "non relational") database provides a mechanism for storage and retrieval of data which is modeled in means other than the tabular relations used in relational databases
- Simplicity of design, simpler “horizontal” scaling, finer control over availability, more flexible data structures



Conclusion

- Offers compromise consistency in favor of availability, partition tolerance and speed
- Concept of eventual consistency where databases changes are propagated to all nodes “eventually”, queries for data might not return updated data immediately or result in reading data that is not accurate



THANKS



@andfanilo



@andfanilo



andfanilo@gmail.com